

early-stopping

November 29, 2023

```
[1]: import tensorflow as tf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import warnings
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.model_selection import train_test_split
import seaborn as sns
```


```
[2]: df=pd.read_csv('/content/data.csv')
df
```

```
[2]:   feature_1  feature_2  label
0   -0.365031 -0.766134      1
1    0.791425  0.424866      1
2   -0.626055  0.529933      1
3    0.719383 -0.598984      0
4    0.875863 -0.796937      0
..      ...      ...      ...
95   0.236494  0.712223      1
96  -0.541747  0.342009      1
97  -0.260651  0.924460      0
98   0.415090 -0.754813      1
99   0.756847  0.115987      1
```

[100 rows x 3 columns]

```
[3]: sns.scatterplot(x='feature_1', y='feature_2', hue='label', data=df)

plt.show()
```



```
[4]: X=df.drop('label',axis=1)
     y=df['label']
```

```
[5]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.10,
     ↪random_state=42)
```


```
[6]: model = Sequential()
     model.add(Dense(256, input_dim=2, activation='relu'))# input_dim=2 vaneko
     ↪features 2 ota xa . 1st layer ma 256 nodes
     model.add(Dense(32, activation='relu')) # second layer ko nodes 32 ota rakheko
     model.add(Dense(1, activation='sigmoid')) #output layer ma euta node
```

```
[7]: model.compile(loss='binary_crossentropy', optimizer='adam',
     ↪metrics=['accuracy'])
```

```
[8]: history = model.fit(X_train, y_train, validation_data=(X_test, y_test),
     ↪epochs=10000, verbose=0)
```

```
[9]: plt.plot(history.history['loss'], label='train')
     plt.plot(history.history['val_loss'], label='test')
     plt.legend()
```

```
plt.show()
```



about 2300 epoch paxi
over fitting va xa

1 EARLY STOPPING

```
[26]: model = Sequential()  
  
model.add(Dense(256, input_dim=2, activation='relu'))  
model.add(Dense(1, activation='sigmoid'))
```

```
[27]: model.compile(loss='binary_crossentropy', optimizer='adam',  
metrics=['accuracy'])
```

```
[28]: callback = EarlyStopping(  
    monitor="val_loss",  
    min_delta=0.0001,  
    patience=200,  
    verbose=1,  
    mode="auto",  
    restore_best_weights=True  
)
```

previous best Val loss - current Val loss

$$0.6926 - 0.6933 = -0.0007$$

```
[29]: history = model.fit(X_train, y_train, validation_data=(X_test, y_test),  
    epochs=10000, callbacks=callback)
```

```
Epoch 1/10000  
3/3 [=====] - 1s 91ms/step - loss: 0.5000 - val_loss: 0.6926 - val_accuracy: 0.5000  
Epoch 2/10000  
3/3 [=====] - 0s 17ms/step - loss: 0.4889 - val_loss: 0.6933 - val_accuracy: 0.5000  
Epoch 3/10000  
3/3 [=====] - 0s 16ms/step - loss: 0.5111 - val_loss: 0.6967 - val_accuracy: 0.5000  
Epoch 4/10000  
3/3 [=====] - 0s 16ms/step - loss: 0.5333 - val_loss: 0.6972 - val_accuracy: 0.5000  
Epoch 5/10000  
3/3 [=====] - 0s 15ms/step - loss: 0.5333 - val_loss: 0.6993 - val_accuracy: 0.3000  
Epoch 6/10000  
3/3 [=====] - 0s 16ms/step - loss: 0.5556 - val_loss: 0.7004 - val_accuracy: 0.3000  
Epoch 7/10000  
3/3 [=====] - 0s 15ms/step - loss: 0.6000 - val_loss: 0.7024 - val_accuracy: 0.4000  
Epoch 8/10000  
3/3 [=====] - 0s 17ms/step - loss: 0.6111 - val_loss: 0.7014 - val_accuracy: 0.4000  
Epoch 9/10000  
3/3 [=====] - 0s 15ms/step - loss: 0.5889 - val_loss: 0.7023 - val_accuracy: 0.4000  
Epoch 10/10000  
3/3 [=====] - 0s 15ms/step - loss: 0.5778 - val_loss: 0.7019 - val_accuracy: 0.4000  
Epoch 11/10000  
3/3 [=====] - 0s 15ms/step - loss: 0.5444 - val_loss: 0.7018 - val_accuracy: 0.4000  
Epoch 12/10000  
3/3 [=====] - 0s 15ms/step - loss: 0.5444 - val_loss: 0.7035 - val_accuracy: 0.4000  
Epoch 13/10000  
3/3 [=====] - 0s 15ms/step - loss: 0.5444 - val_loss: 0.7050 - val_accuracy: 0.4000  
Epoch 14/10000  
3/3 [=====] - 0s 16ms/step - loss: 0.5444 - val_loss: 0.7062 - val_accuracy: 0.4000  
Epoch 15/10000  
3/3 [=====] - 0s 17ms/step - loss: 0.5444 - val_loss: 0.7062 - val_accuracy: 0.4000
```

**Early stoping is triggered
and now it waits for 200
epoches more to get less val
loss than 0.6926 until the
diff of previous best
vloss(0.6926) - curent
weight will be pos >= 0.0001**

$$0.3081 - 0.3124 = - 0.0043$$

```

0.8444 - val_loss: 0.3235 - val_accuracy: 0.8000
Epoch 2016/10000
3/3 [=====] - 0s 33ms/step - loss: 0.3601 - accuracy:
0.8444 - val_loss: 0.3243 - val_accuracy: 0.8000
Epoch 2017/10000
3/3 [=====] - 0s 31ms/step - loss: 0.3602 - accuracy:
0.8333 - val_loss: 0.3211 - val_accuracy: 0.8000
Epoch 2018/10000
3/3 [=====] - 0s 29ms/step - loss: 0.3601 - accuracy:
0.8333 - val_loss: 0.3205 - val_accuracy: 0.8000
Epoch 2019/10000
3/3 [=====] - 0s 25ms/step - loss: 0.3603 - accuracy:
0.8333 - val_loss: 0.3244 - val_accuracy: 0.9000
Epoch 2020/10000
3/3 [=====] - 0s 25ms/step - loss: 0.3610 - accuracy:
0.8444 - val_loss: 0.3266 - val_accuracy: 0.8000
Epoch 2021/10000
3/3 [=====] - 0s 31ms/step - loss: 0.3606 - accuracy:
0.8333 - val_loss: 0.3251 - val_accuracy: 0.8000
Epoch 2022/10000
3/3 [=====] - 0s 31ms/step - loss: 0.3600 - accuracy:
0.8333 - val_loss: 0.3252 - val_accuracy: 0.8000
Epoch 2023/10000
3/3 [=====] - 0s 30ms/step - loss: 0.3602 - accuracy:
0.8333 - val_loss: 0.3249 - val_accuracy: 0.8000
Epoch 2024/10000
3/3 [=====] - 0s 31ms/step - loss: 0.3597 - accuracy:
0.8333 - val_loss: 0.3221 - val_accuracy: 0.8000
Epoch 2025/10000
3/3 [=====] - 0s 42ms/step - loss: 0.3598 - accuracy:
0.8333 - val_loss: 0.3170 - val_accuracy: 0.8000
Epoch 2026/10000
3/3 [=====] - 0s 32ms/step - loss: 0.3598 - accuracy:
0.8444 - val_loss: 0.3139 - val_accuracy: 0.8000
Epoch 2027/10000
3/3 [=====] - 0s 31ms/step - loss: 0.3597 - accuracy:
0.8444 - val_loss: 0.3087 - val_accuracy: 0.8000
Epoch 2028/10000
3/3 [=====] - 0s 31ms/step - loss: 0.3602 - accuracy:
0.8444 - val_loss: 0.3124 - val_accuracy: 0.8000
Epoch 2029/10000
3/3 [=====] - 0s 35ms/step - loss: 0.3603 - accuracy:
0.8333 - val_loss: 0.3225 - val_accuracy: 0.8000
Epoch 2031/10000
3/3 [=====] - 0s 31ms/step - loss: 0.3601 - accuracy:

```

early stoping trigger who aba next
 200 epochs samma yo loss vanda
 kunai kam loss aauxa ki aaudina
 herxa ra yo (0.3087 - current
 weights) ko values ≥ 0.001 aauxa
 ki nai herxa aayena vane traing
 stop hunca

2027 epoch paxi next 200 epoch
samma ma kunai epoch ma val
loss less than 0.3087 aayena.
sabai ko previous best
weight(0.3087) - current weight
garda < 0.001 nai aayo (neg
value) so training rokiyo

```
0.8556 - val_loss: 0.3239 - val_accuracy: 0.8000
Epoch 2224/10000
3/3 [=====] - 0s 20ms/step - loss: 0.3593 - accuracy: 0.8000
0.8556 - val_loss: 0.3183 - val_accuracy: 0.8000
Epoch 2225/10000
3/3 [=====] - 0s 28ms/step - loss: 0.3593 - accuracy: 0.8000
0.8444 - val_loss: 0.3132 - val_accuracy: 0.8000
Epoch 2226/10000
3/3 [=====] - 0s 29ms/step - loss: 0.3593 - accuracy: 0.8000
0.8444 - val_loss: 0.3154 - val_accuracy: 0.8000
Epoch 2227/10000
1/3 [=====>...] - ETA: 0s - loss: 0.3576 - accuracy: 0.8438
Restoring model weights from the end of the best epoch: 2027.
3/3 [=====] - 0s 34ms/step - loss: 0.3592 - accuracy: 0.8444
0.8444 - val_loss: 0.3207 - val_accuracy: 0.8000
Epoch 2227: early stopping
```

```
[30]: plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='test')
plt.legend()
plt.show()
```

